

Database Management Systems INFO 210

The Relational Model Lecture 5

Franz Wotawa

TU Graz, Institut for Software Technologie
Inffeldgasse 16b/2
wotawa@ist.tugraz.at

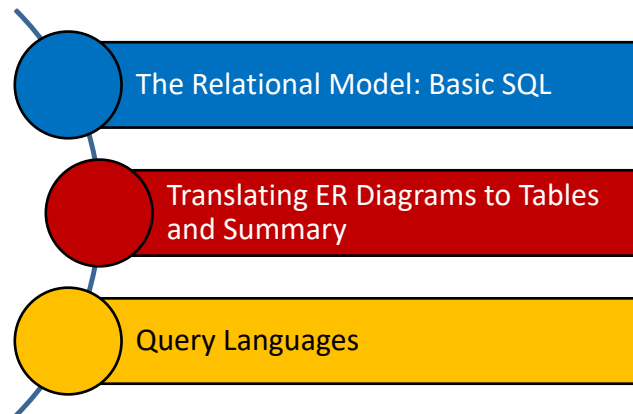
1

Today...

- **Last Session:**
 - The ER model
- **Today's Session:**
 - The relational model
 - Basic SQL
 - ER to relational databases

2

Outline



3

What is the Relational Model?

- The relational model adopts a “tabular” representation
 - A database is a *collection* of one or more [relations](#)
 - Each relation is a *table* with rows and columns
- What is unique about the relational model as opposed to older data models?
 - Its simple data representation
 - Ease with which complex queries can be expressed

4

Basic Constructs

- The main construct in the relational model is the *relation*
- A relation consists of:
 1. A *schema* which includes:
 - The relation's name
 - The name of each column
 - The *domain* of each column
 2. An *instance* which is a set of tuples
 - Each tuple has the same number of columns as the relation schema

5

Creating Relations in SQL

- S1 can be used to create the "Students" relation
- S2 can be used to create the "Enrolled" relation

```
CREATE TABLE Students
(sid CHAR(20),
 name CHAR(20),
 login CHAR(10),
 age INTEGER,
 gpa REAL);
```

S1

```
CREATE TABLE Enrolled
(sid CHAR(20),
 cid CHAR(20),
 grade CHAR(2));
```

S2

The DBMS enforces domain constraints whenever tuples are added or modified

6

Adding and Deleting Tuples

- We can insert a single tuple to the “Students” relation using:

```
INSERT INTO Students (sid, name, login, age, gpa)
VALUES ('53688', 'Smith', 'smith@ee', 18, 3.2);
```

- We can delete all tuples from the “Students” relation which satisfy some condition (e.g., name = Smith):

```
DELETE
FROM Students S
WHERE S.name = 'Smith';
```

Powerful variants of these commands are available; *more on this next week!*

7


Integrity Constraints (ICs)

- An **IC** is a condition that must be true for *any* instance of the database (e.g., *domain constraints*)
 - ICs are specified when schemas are defined
 - ICs are *checked* when relations are modified
- A **legal** instance of a relation is one that satisfies all specified ICs
 - DBMS should not allow illegal instances
- If the DBMS checks ICs, stored data is more faithful to real-world meaning
 - Avoids data entry errors, too!

8

Keys

- Keys help associate tuples in different relations
- Keys are one form of integrity constraints (ICs)



9

Primary and Candidate Keys in SQL

- Many candidate keys (specified using **UNIQUE**) can be designated and one is chosen as a *primary key*
- Keys must be used carefully!
- “For a given student and course, there is a single grade”

```
CREATE TABLE Enrolled
(sid CHAR(20)
 cid CHAR(20),
 grade CHAR(2),
 PRIMARY KEY (sid,cid));
```

vs.

```
CREATE TABLE Enrolled
(sid CHAR(20)
 cid CHAR(20),
 grade CHAR(2),
 PRIMARY KEY (sid),
 UNIQUE (cid, grade));
```

Q: What does this mean?

10

Primary and Candidate Keys in SQL

- Many candidate keys (specified using **UNIQUE**) can be designated and one is chosen as a *primary key*
- Keys must be used carefully!
- “For a given student and course, there is a single grade”

```
CREATE TABLE Enrolled
(sid CHAR(20)
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid));
```

vs.

```
CREATE TABLE Enrolled
(sid CHAR(20)
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid),
UNIQUE (cid, grade));
```

“A student can take only one course, and no two students in a course receive the same grade”

11

Foreign Keys in SQL

- Example: Only existing students may enroll for courses

```
CREATE TABLE Enrolled
(sid CHAR(20),cid CHAR(20),grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid) REFERENCES students );
```

Enrolled			Students				
sid	cid	grade	sid	name	login	age	gpa
53666	15-101	C	53666	Jones	jones@cs	18	3.4
53666	18-203	B	53688	Smith	smith@cs	18	3.2
53650	15-112	A	53650	Smith	smith@math	19	3.8
53666	15-105	B					

12

Enforcing Referential Integrity

- What should be done if an “Enrolled” tuple with a non-existent student id is inserted? (*Reject it!*)
- What should be done if a “Students” tuple is deleted?
 - Disallow its deletion
 - Delete all Enrolled tuples that refer to it
 - Set sid in Enrolled tuples that refer to it to a *default sid*
 - Set sid in Enrolled tuples that refer to it to a special value *null*, denoting ‘unknown’ or ‘inapplicable’
- What if a “Students” tuple is updated?

13

Referential Integrity in SQL

- SQL/92 and SQL:1999 support all 4 options on deletes and updates
 - Default is **NO ACTION** (i.e., *delete/update is rejected*)
 - **CASCADE** (also delete all tuples that refer to the deleted tuple)
 - **SET NULL / SET DEFAULT** (sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled
(sid CHAR(20),
 cid CHAR(20),
 grade CHAR(2),
 PRIMARY KEY (sid,cid),
 FOREIGN KEY (sid)
 REFERENCES Students
 ON DELETE CASCADE
 ON UPDATE SET DEFAULT );
```

What does this mean?

14

Views

- A **view** is a table whose rows are not explicitly stored but computed as needed

```
CREATE VIEW YoungActiveStudents (name, grade)
AS SELECT S.name, E.grade
FROM Students S, Enrolled E
WHERE S.sid = E.sid and S.age < 21;
```

- Views can be queried
 - Querying YoungActiveStudents would necessitate computing it first then applying the query on the result as being like any other relation
- Views can be dropped using the **DROP VIEW** command
 - How to handle **DROP TABLE** if there's a view on the table?
 - DROP TABLE command has options to let the user specify this

15

Views and Security

- Views can be used to present necessary information, while hiding details in underlying relation(s)
 - If the schema of an old relation is *changed*, a view can be defined to represent the old schema
 - This allows applications to *transparently* assume the old schema
- Views can be defined to give a group of users access to just the information they are allowed to see
 - E.g., we can define a view that allows students to see other students' names and ages, but not GPAs (also students can be prevented from accessing the underlying "Students" relation)

16

Views and Security

- Views can be used to present necessary information, while hiding details in underlying relation(s)
 - If the schema of an old relation is *changed*, a view can be defined to represent the old schema
 - This allows applications to *transparently* assume the old schema
- Views can be defined to give a group of users access to just the information they are allowed to see
 - E.g., we can define a view that allows students to see other students' names and ages, but not GPAs (also students can be prevented from accessing the underlying "Students" relation)

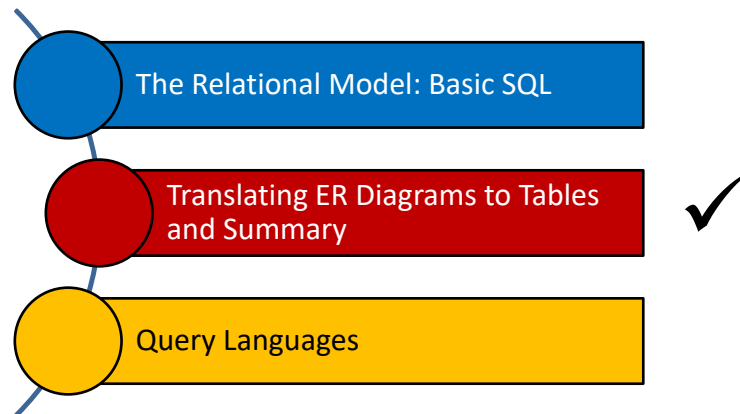
17

Views and Security

- Views can be used to present necessary information, while hiding details in underlying relation(s)
 - If the schema of an old relation is *changed*, a view can be defined to represent the old schema
 - This allows applications to *transparently* assume the old schema
- Views can be defined to give a group of users access to just the information they are allowed to see
 - E.g., we can define a view that allows students to see other students' names and ages, but not GPAs (also students can be prevented from accessing the underlying "Students" relation)


18

Outline



19

Strong Entity Sets to Tables




20

Relationship Sets to Tables

- In translating a relationship set to a relation, attributes of the relation must include:
 1. Keys for each participating entity set (as foreign keys)
 - This set of attributes forms a *superkey* for the relation
 2. All descriptive attributes
- Relationship sets
 - 1-to-1, 1-to-many, and many-to-many
 - Key/Total/Partial participation

21


M-to-N Relationship Sets to Tables



```
CREATE TABLE Works_In(
  ssn CHAR(11),
  did INTEGER,
  since DATE,
  PRIMARY KEY (ssn, did),
  FOREIGN KEY (ssn)
    REFERENCES Employees,
  FOREIGN KEY (did)
    REFERENCES Departments);
```

22

1-to-M Relationship Sets to Tables



```

CREATE TABLE Manages(
  ssn CHAR(11),
  did INTEGER,
  since DATE,

  PRIMARY KEY (did),
  FOREIGN KEY (ssn)
  REFERENCES Employees,
  FOREIGN KEY (did)
  REFERENCES Departments);
  
```


```

CREATE TABLE Departments(
  did INTEGER,
  dname CHAR(20),
  budget REAL,
  PRIMARY KEY (did)
);
  
```

Approach 1:
Create separate tables for Manages and Departments

23

1-to-M Relationship Sets to Tables



```

CREATE TABLE Dept_Mgr(
  ssn CHAR(11),
  did INTEGER,
  since DATE,
  dname CHAR(20),
  budget REAL,
  PRIMARY KEY (did),
  FOREIGN KEY (ssn)
  REFERENCES Employees);
  
```

Can ssn take a
null value?

Approach 2:
Create a table for only the Departments entity set (i.e., take advantage of the key constraint)

24


One-Table vs. Two-Table Approaches

- The **one-table approach**:
 - (+) Eliminates the need for a separate table for the involved relationship set (e.g., Manages)
 - (+) Queries can be answered without combining information from two relations
 - (-) Space could be wasted!
 - What if several departments have no managers?
- The **two-table approach**:
 - The opposite of the one-table approach!

25

Translating Relationship Sets with Participation Constraints

- What does the following ER diagram entail (with respect to Departments and Managers)?



Every *did* value in Departments table must appear in a row of the Manages table- *if defined*- (with a non-null *ssn* value!)

26

Translating Relationship Sets with Participation Constraints

- Here is how to create the “Dept_Mgr” table using the one-table approach:

```
CREATE TABLE Dept_Mgr(
  did INTEGER,
  dname CHAR(20),
  budget REAL,
  ssn CHAR(11) NOT NULL,
  since DATE,
  PRIMARY KEY (did),
  FOREIGN KEY (ssn) REFERENCES Employees,
  ON DELETE NO ACTION);
```

Can this be captured using the two-table approach?

27

Translating Relationship Sets with Participation Constraints

- Here is how to create the “Dept_Mgr” table using the one-table approach:


```
CREATE TABLE Dept_Mgr(
  did INTEGER,
  dname CHAR(20),
  budget REAL,
  ssn CHAR(11) NOT NULL,
  since DATE,
  PRIMARY KEY (did),
  FOREIGN KEY (ssn) REFERENCES Employees,
  ON DELETE SET NULL);
```

Would this work?

28

Translating Weak Entity Sets

- A weak entity set always:
 - Participates in a one-to-many binary relationship
 - Has a key constraint and total participation




- Which approach is ideal for that?
 - The one-table approach

29

Translating Weak Entity Sets

- Here is how to create “Dep_Policy” using the one-table approach

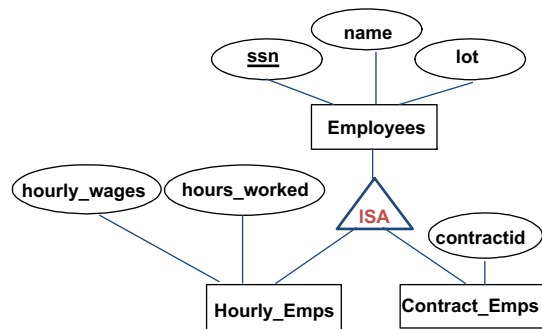


```
CREATE TABLE Dep_Policy (
  dname CHAR(20),
  age    INTEGER,
  cost   REAL,
  ssn    CHAR(11) NOT NULL,
  PRIMARY KEY (dname, ssn),
  FOREIGN KEY (ssn) REFERENCES Employees,
  ON DELETE CASCADE);
```

30

Translating ISA Hierarchies to Relations

- Consider the following example:



31

Translating ISA Hierarchies to Relations

- General approach:
 - Create 3 relations: "Employees", "Hourly_Emps" and "Contract_Emps"

EMP (ssn, name, lot)


Blue		
Black		
Green		
Red		

H_EMP(ssn, h_wg, h_wk)

Blue		
Red		

CONTR(ssn, cid)

Green	
-------	--

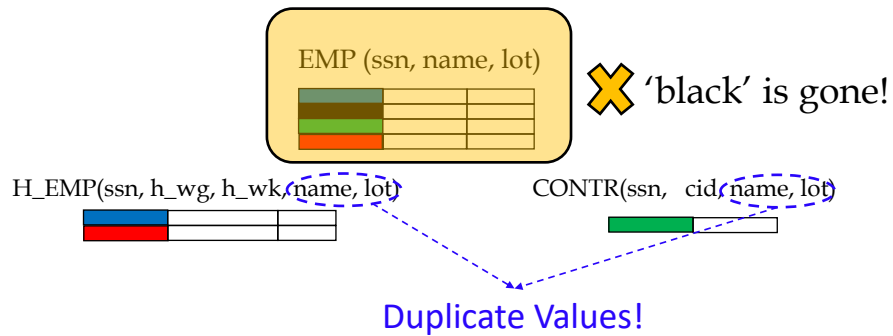



- How many times do we record an employee?
- What to do on deletions?
- How to retrieve **all** info about an employee?

32

Translating ISA Hierarchies to Relations

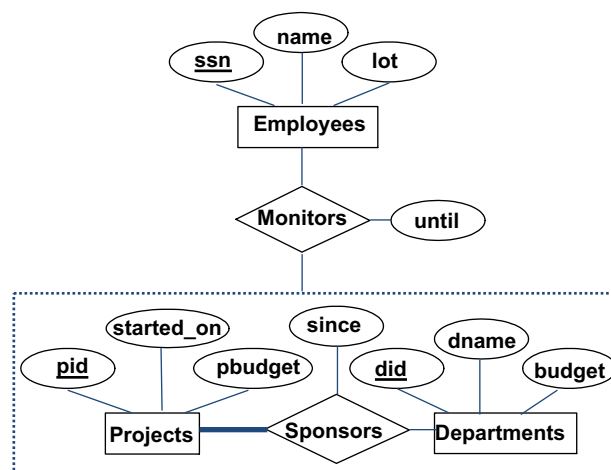
- Alternatively:
 - Just create 2 relations “Hourly_Emps” and “Contract_Emps”
 - Each employee **must be** in one of these two subclasses



33

Translating Aggregations

- Consider the following example:

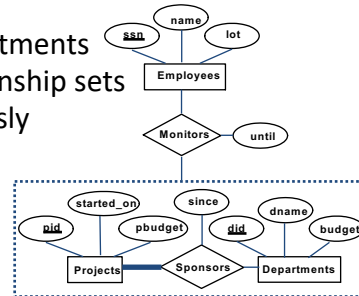


34

Translating Aggregations

- Standard approach:

- The Employees, Projects and Departments entity sets and the Sponsors relationship sets are translated as described previously



- For the Monitors relationship, we create a relation with the following attributes:

- The key attribute of Employees (i.e., ssn)
 - The key attributes of Sponsors (i.e., did, pid)
 - The descriptive attributes of Monitors (i.e., until)

35

The Relational Model: A Summary

- A tabular representation of data
- Simple and intuitive, currently one of the most widely used
 - Object-Relational Mapping (ORM) hides the relational model
 - Non-relational NoSQL model is gaining ground
- Integrity constraints can be specified (by the DBA) based on application semantics (DBMS checks for violations)
 - Two important ICs: primary and foreign keys
 - Also: not null, unique
 - In addition, we *always* have domain constraints
- Mapping from ER to Relational is (fairly) straightforward!

36

ER to Tables - Summary of Basics

- Strong entities:
 - Key -> primary key
- (Binary) relationships:
 - Get keys from all participating entities:
 - 1:1 -> either key can be the primary key
 - 1:N -> the key of the 'N' part will be the primary key
 - M:N -> both keys will be the primary key
- Weak entities:
 - Strong key + partial key -> primary key
 - ON DELETE CASCADE

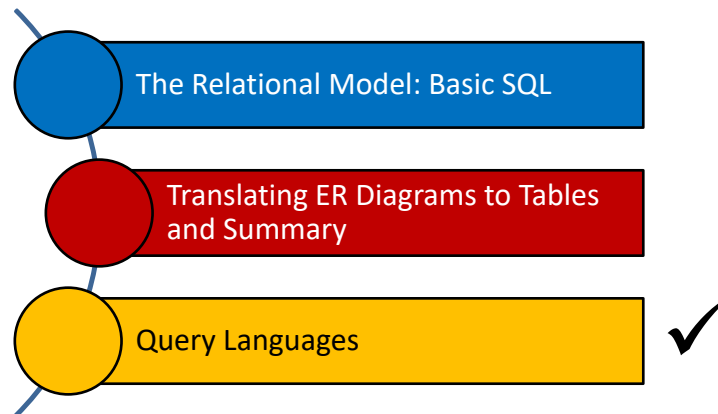
37

ER to Tables - Summary of Advanced

- Total/Partial participation:
 - NOT NULL
- Ternary relationships:
 - Get keys from all; decide which one(s) -> primary Key
- Aggregation: like relationships
- ISA:
 - 3 tables (most general)
 - 2 tables ('total coverage')

38

Outline



39

Relational Query Languages

- **Query languages** (QLs) allow *manipulating* and *retrieving* data from databases
- The relational model supports simple and powerful QLs:
 - Strong formal foundation based on logic
 - High amenability for effective optimizations
- **Query Languages != programming languages!**
 - QLs are not expected to be “Turing complete”
 - QLs are not intended to be used for complex calculations
 - QLs support easy and efficient access to large datasets

40

Formal Relational Query Languages

- There are two mathematical Query Languages which form the basis for commercial languages (e.g., SQL)
 - **Relational Algebra**
 - Queries are composed of operators
 - Each query describes a step-by-step procedure for computing the desired answer
 - Very useful for representing *execution plans*
 - **Relational Calculus**
 - Queries are subsets of first-order logic
 - Queries describe desired answers without specifying how they will be computed
 - A type of *non-procedural* (or *declarative*) formal query language

41

Next Class

Relational Algebra

42